

Multilevel Analysis of Software Systems


Naser Ezzati

Brock University

www.github.com/naser/ubrock

Goal: Debug Software in Production

2



Debuggers and Profilers

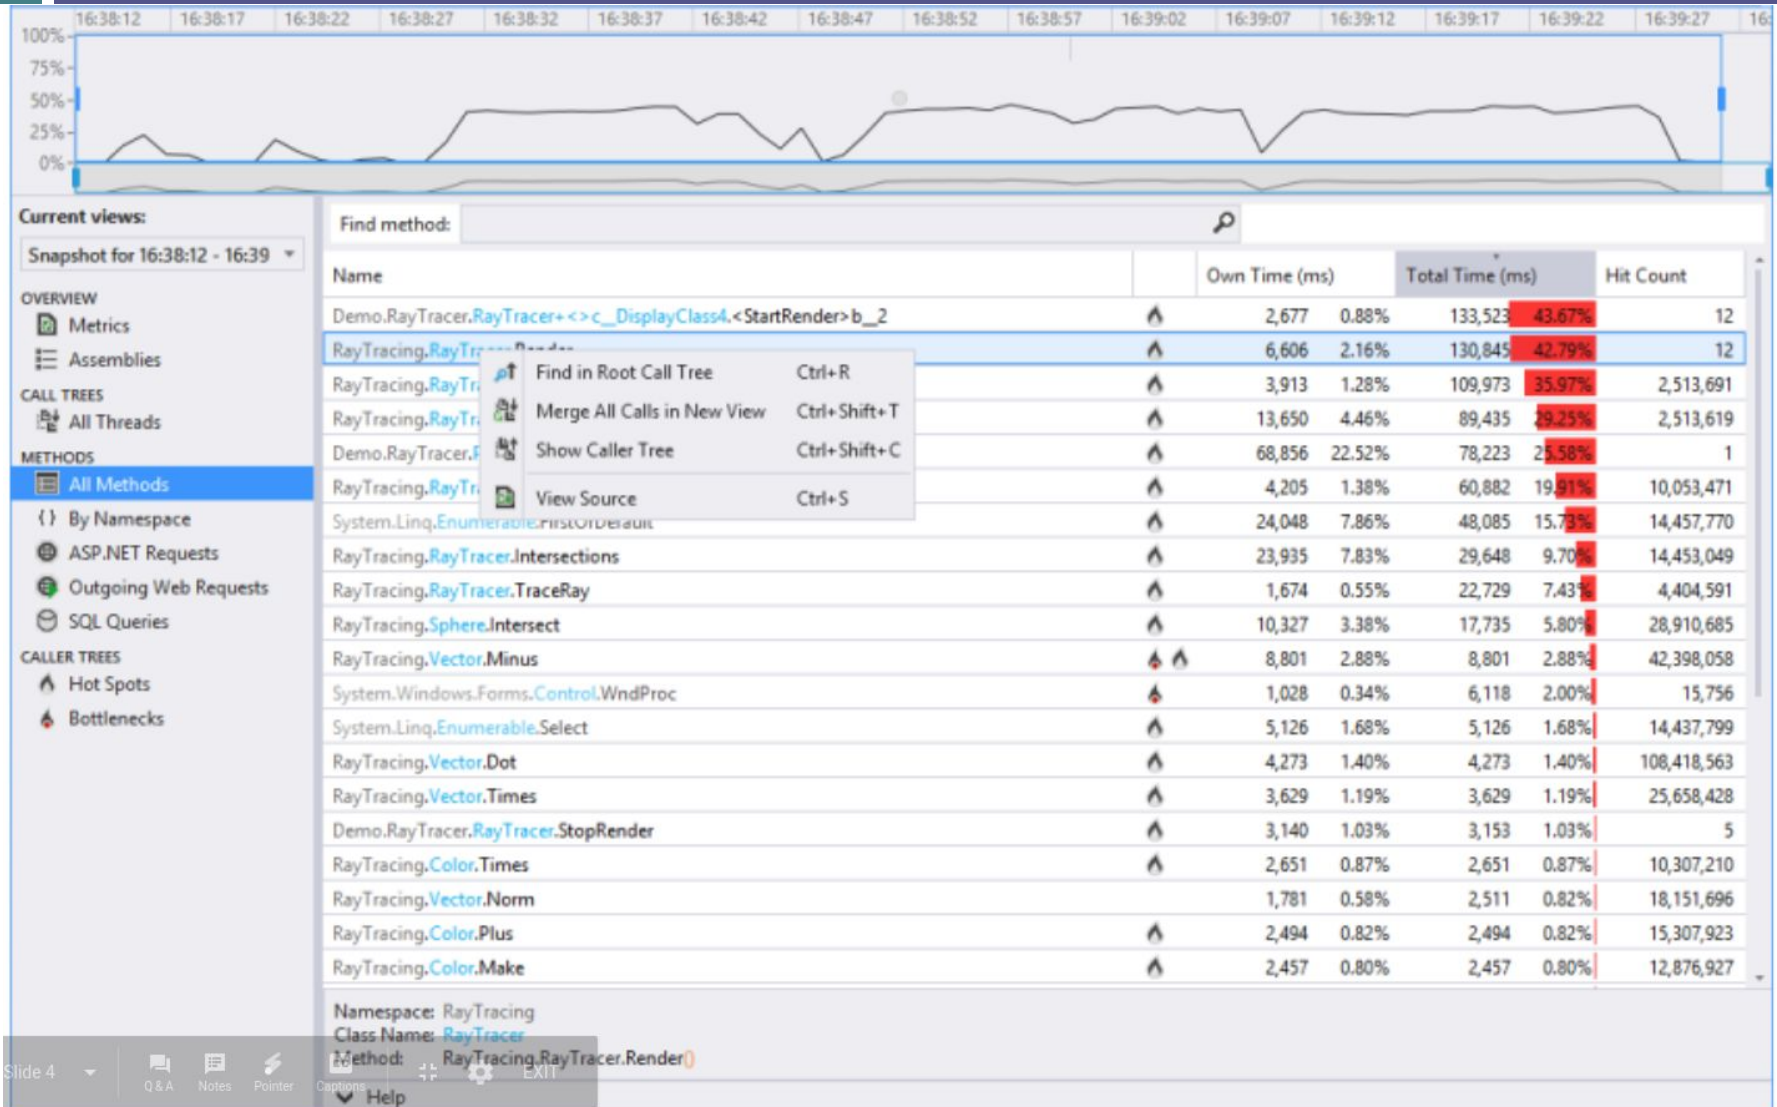
3

```
##### # # # # (gdb) break main
# # ## # # # Breakpoint 1 at 0x8048426: file hello10.c, line 6.
# # # # # (gdb) run
# ##### # # # # Starting program: /home/gary/hello10
# # # # #
# # # # # Breakpoint 1, main () at hello10.c:6
# # # # # 6 for(i=0;i<10;i++)
##### # # #####
(gdb) █

#####
# # ##### # # # # # # # # # #
# # # # # # # # # # # # # # #
# # ##### # # # # # # # #
# # # # # # # # # # # # # #
# # # # # # # # # # # # # #
##### ##### # # # # # # # # # #
```

Profilers

4



Tracing

5



Instrumentation

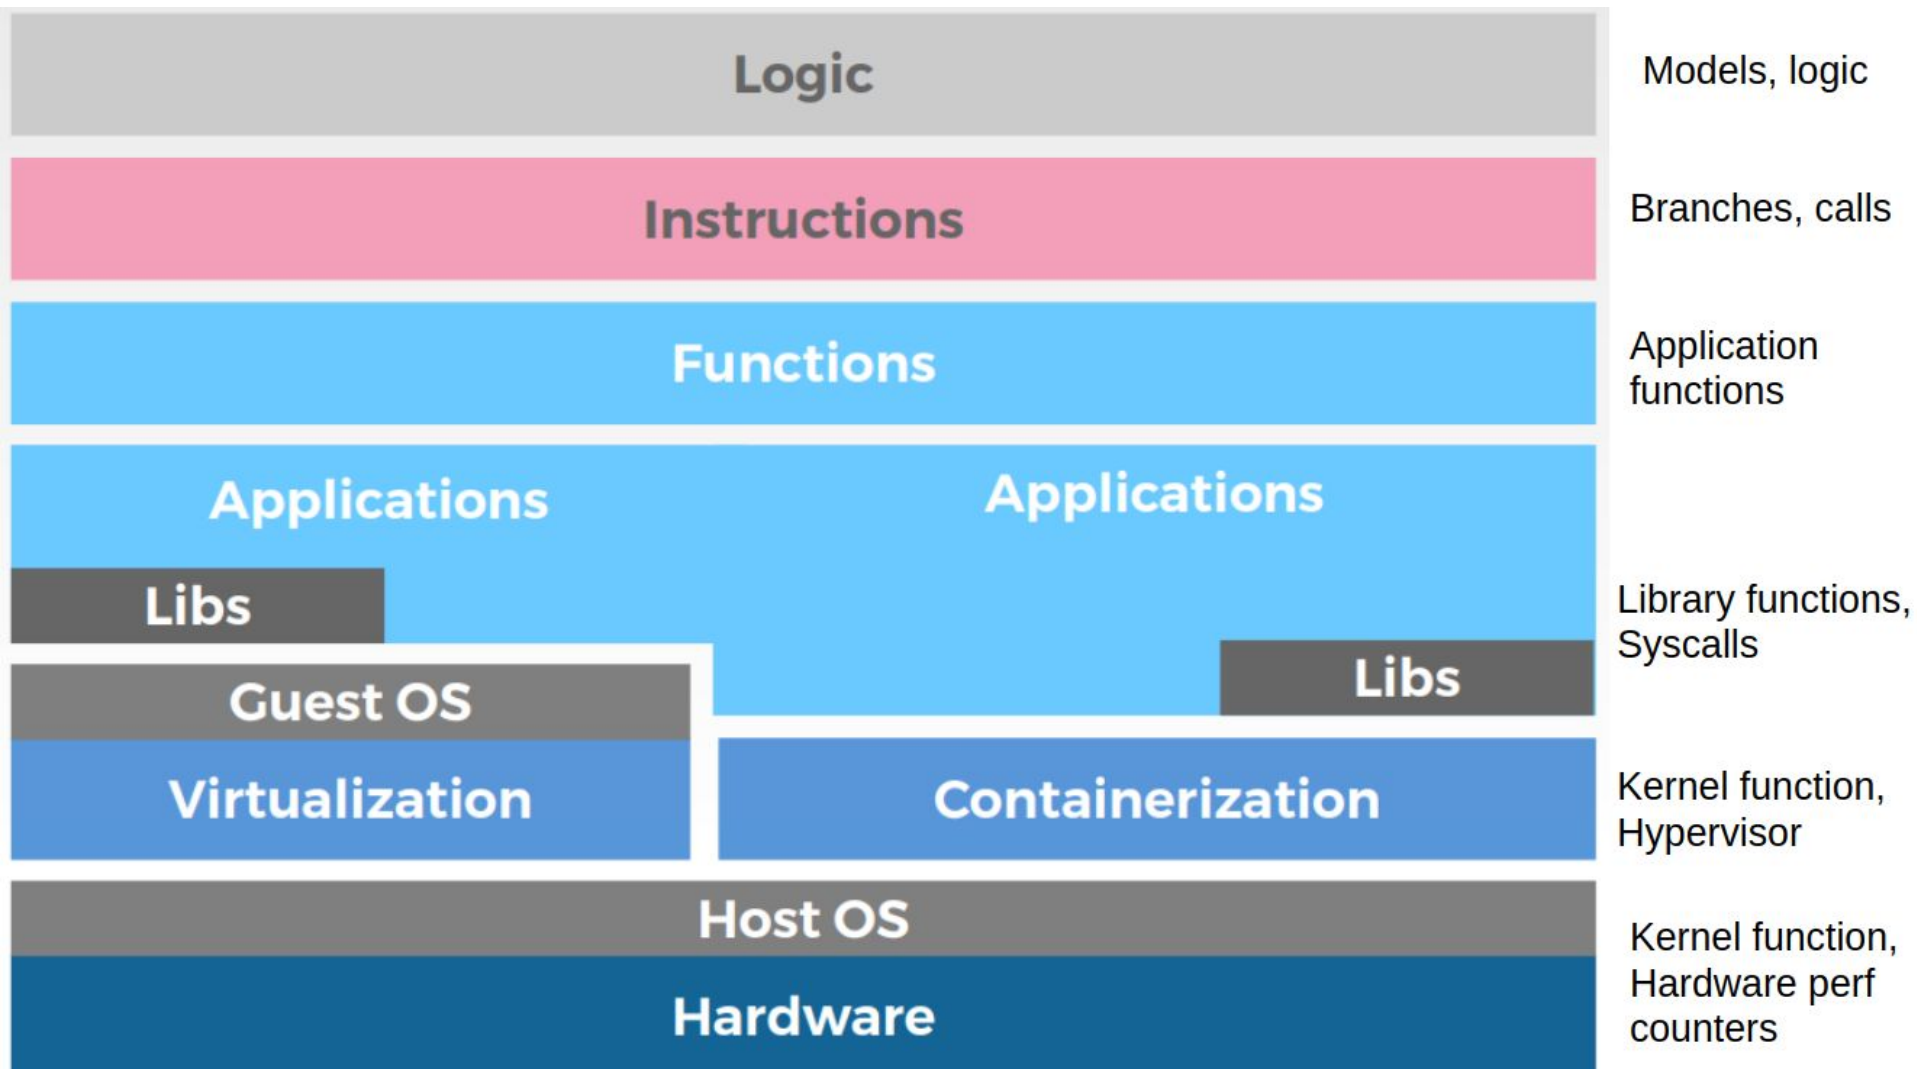
6

```
int my_func(void* my_value) {  
    int i, ret;  
  
    printf("%s: Entering my function", timestamp) //<--  
    for (i = 0; i < MAX; i++) {  
        ret = do_something_with(my_value, i);  
        printf("In for loop %d", i); //<--  
    }  
    printf("%s: Done: %d", timestamp, ret); //<--  
  
    return ret;  
}
```

```
$ ./myprog  
12345674:Entering my function // <-- event  
12345675:In for loop 0 // <-- event  
12345676:In for loop 1 // <-- event  
12345677:In for loop 2 // <-- event  
12345678:Done: 10 // <-- event
```


Multilevel Tracing

7




Why Multilevel Analysis?

8




Why Multilevel Analysis?

9



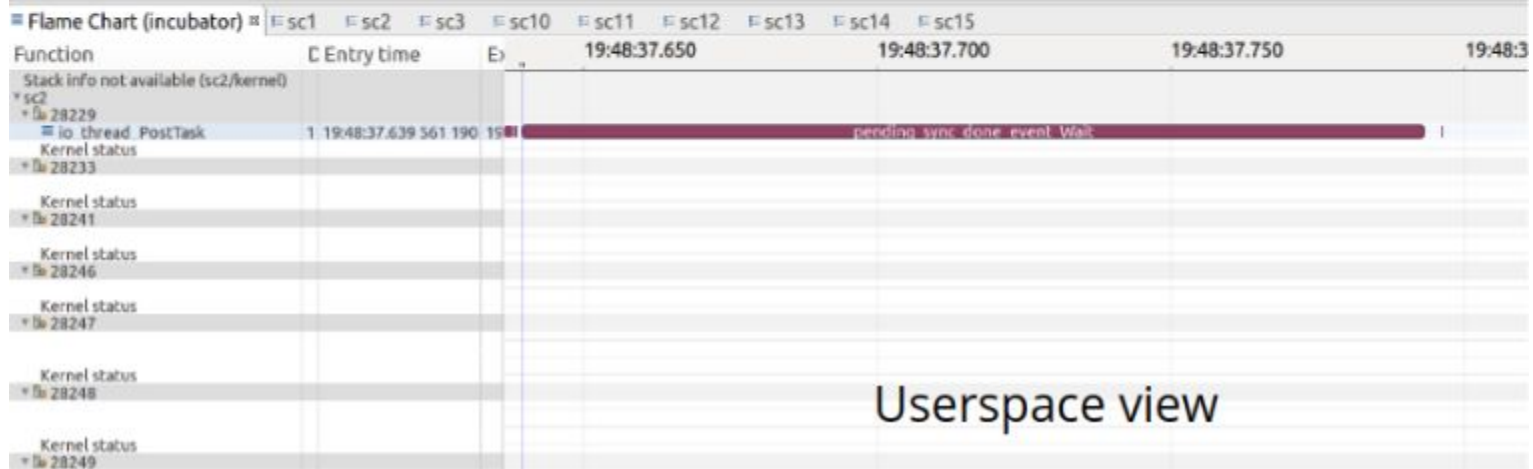
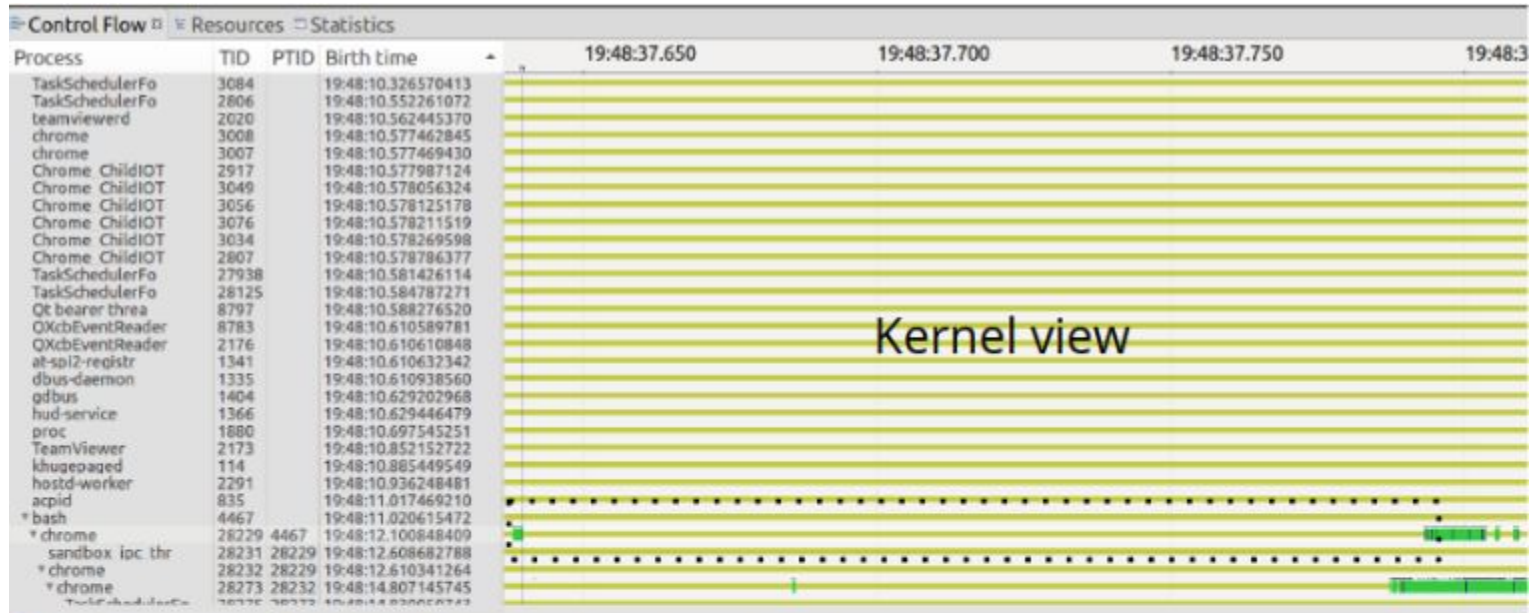
Why Multilevel Analysis?

10



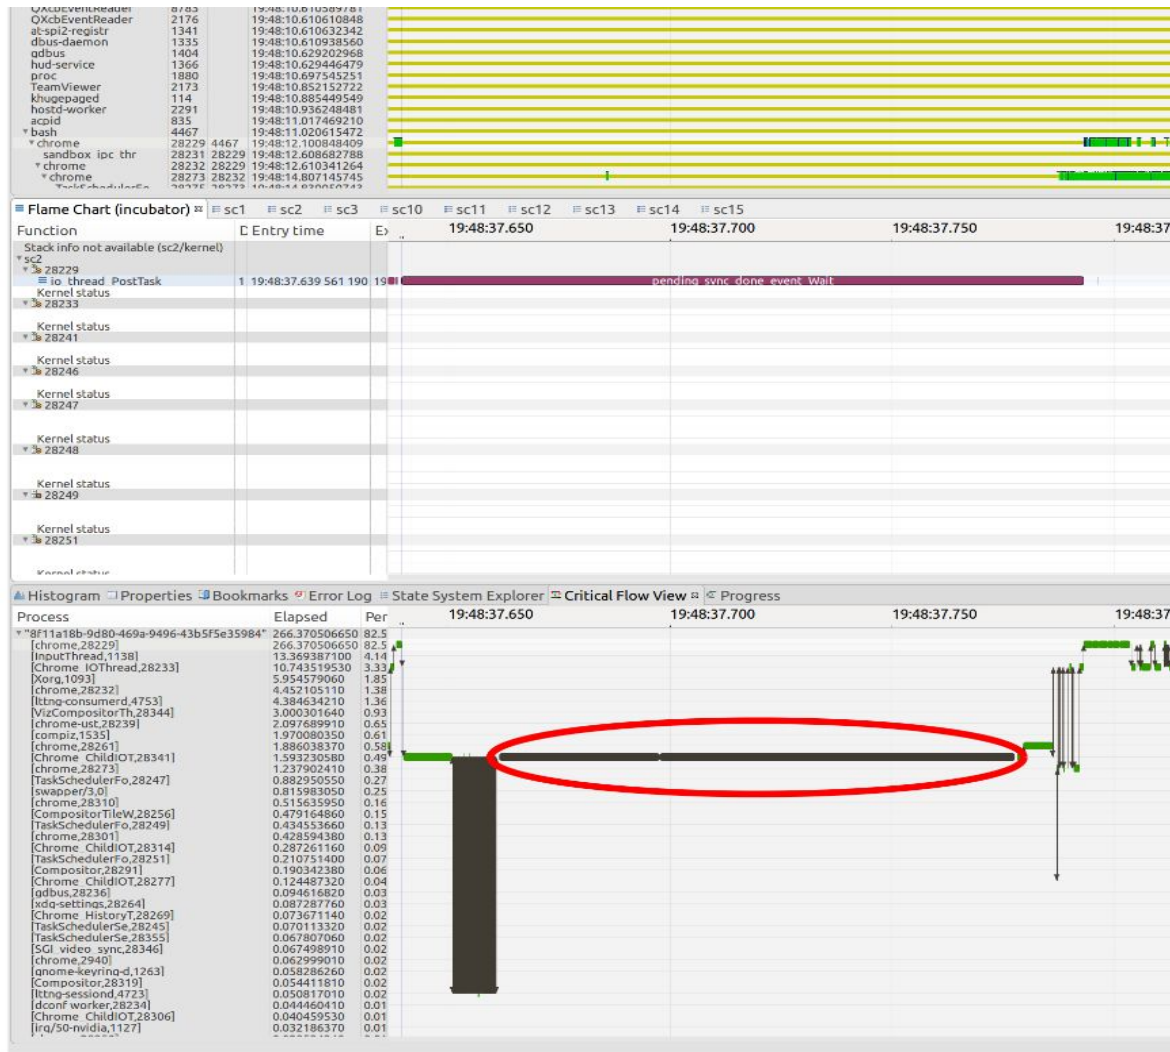
Why Multilevel Analysis?

11




Why Multilevel Analysis?

12




Web Requests Latency Problem

13



Manual Investigation?

14



Trace View

15

- Huge amount of information not related to the observed problem.



Behavioral Analysis

16

- Can we **cluster** executions/requests based on their **behavior**?
- Can we **identify the differences** between groups of requests?
 - Can we find similar requests within a cluster?
 - Can we find different requests between clusters?



Converting Trace to State Sequences

17



ap_R → ap_R → my_U → my_R → ap_P → ap_R → ap_R → ...

time: 10ms

ap_R → ap_R → my_U → my_R → ap_P → ap_R → my-R → ...

time: 14ms

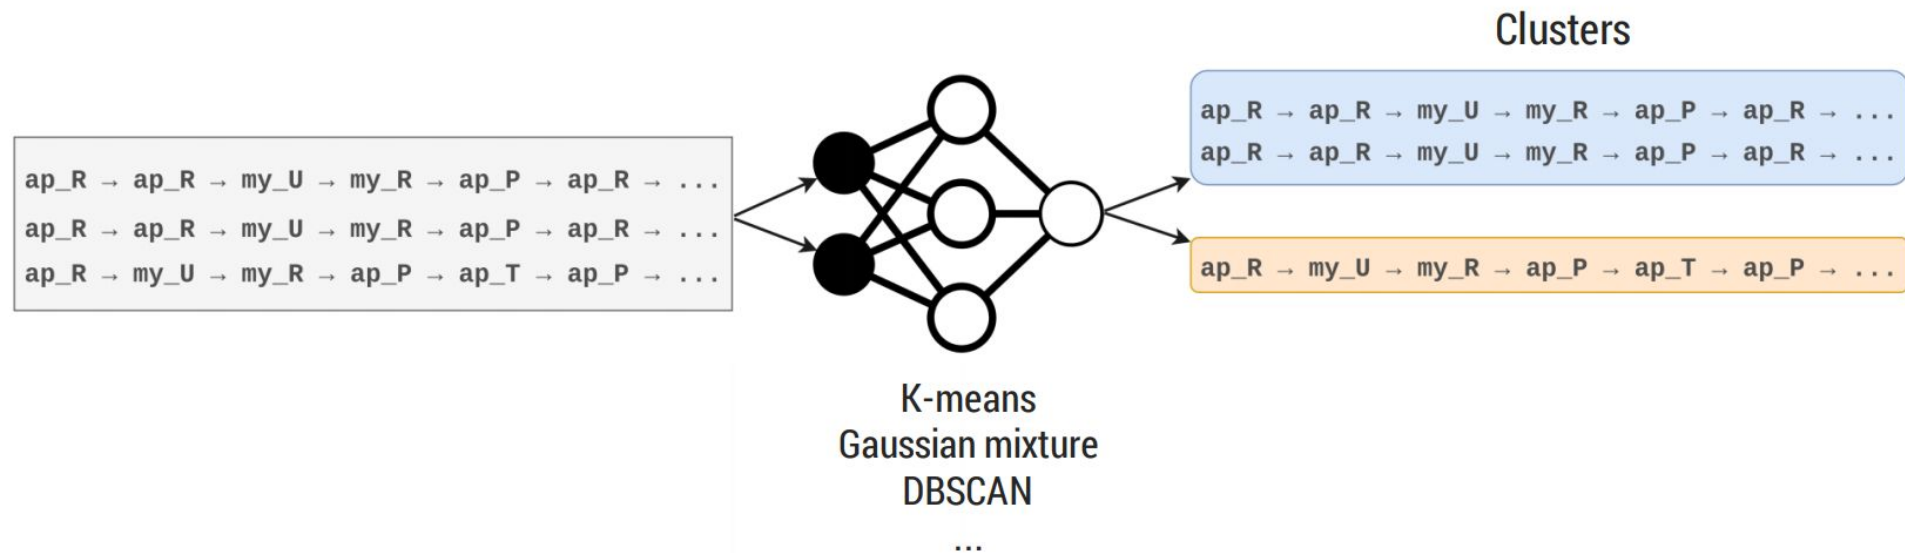
ap_R → my_U → my_R → ap_P → ap_T → ap_P → ap-R → ...

time: 36ms

...


Clustering

18



Resource Usages Clustering

19

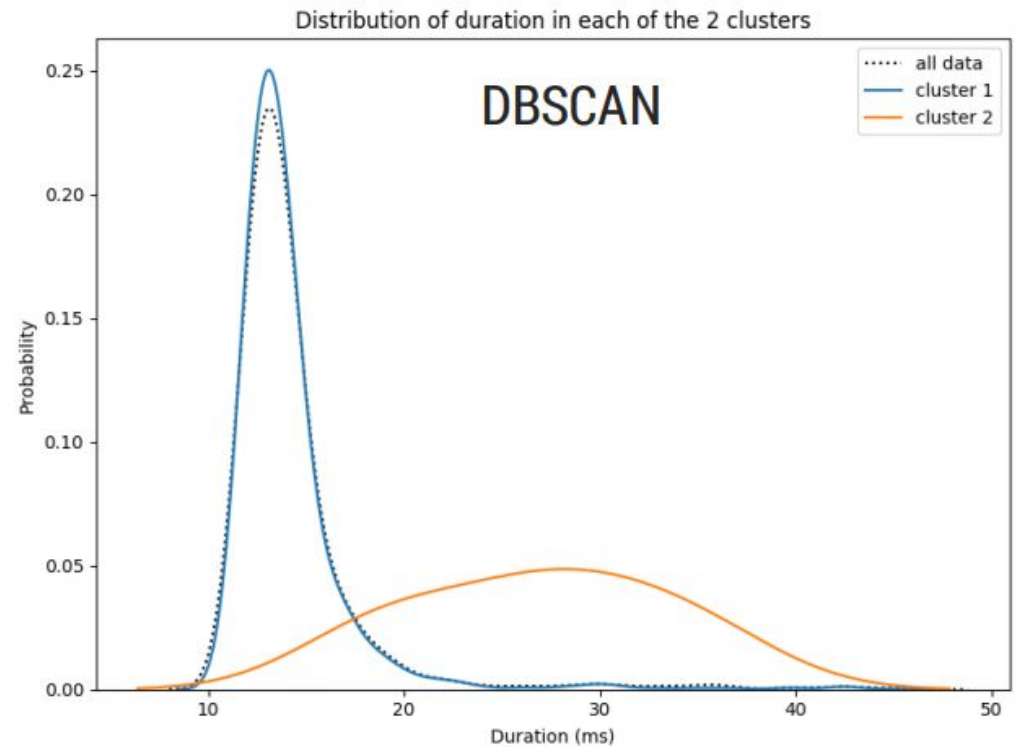
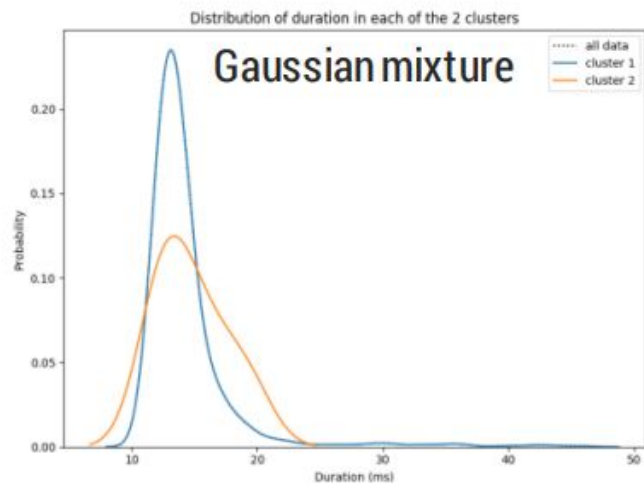



Metrics:

Wait AVG for Disk average time
Wait AVG for Net average time
Wait AVG for Timer average time
Wait AVG for Task average time
Wait AVG for CPU average time
The frequency of wait for Disk
The frequency of wait for Net
The frequency of wait for Timer
The frequency of wait for Task
The frequency of wait for CPU
The frequency of read from Disk
The frequency of write to Disk
Total Block Disk Read
Total Latency Disk Read
Total Block Disk Write
Total Latency Disk Write
Disk Cache Hit/Miss Ratio


Clustering Result

20



DBScan Clustering

21



Analyse the Clustering Result

22


(C0) A latency of 10 ms. Most requests belong to this category

(C1) A latency of [180 ms, 220 ms]

(C2) A latency of [600 ms, 800 ms]

C0: No disk!

C1: and C2 partially/totally served from disk!




Analyse the Clustering Result

23

- (C0) A latency of 10 ms. Most requests belong to this category
- { (C1) A latency of [180 ms, 220 ms]
- { (C2) A latency of [600 ms, 800 ms]


Cache Hit Feature helps to understand the origin of slowness of some requests, but

Why C2 requests take more time to read from the disk, while they have the same cache hit/miss ratio?




Analyse the Clustering Result

24



C2 requests are processed when the disk drive is busy processing other requests from other processes.





- A backup.sh process
- Other web requests



Automatic Investigation

25

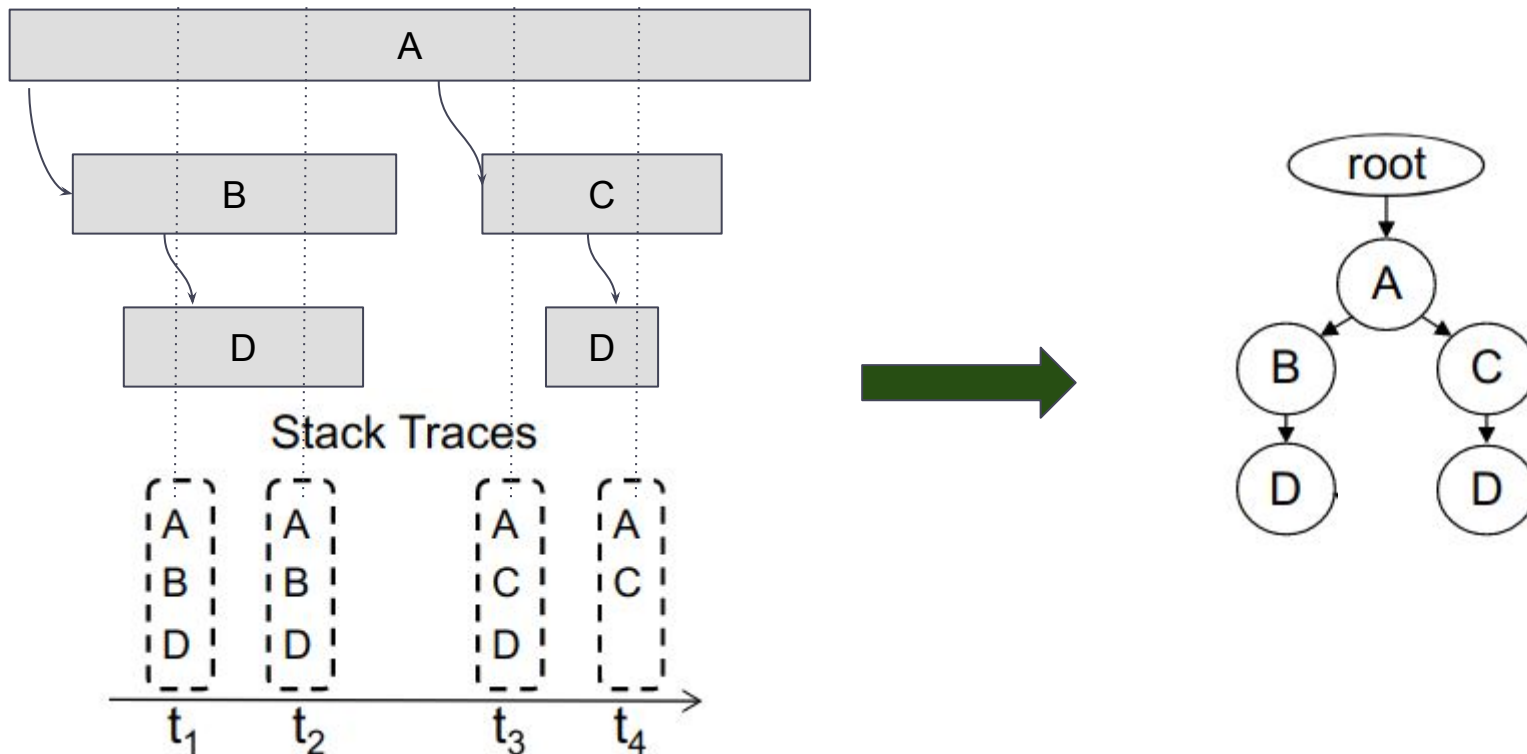
- Extract CCT for each request



Automatic Investigation

26


- Extract CCT for each request



Automatic Investigation

27

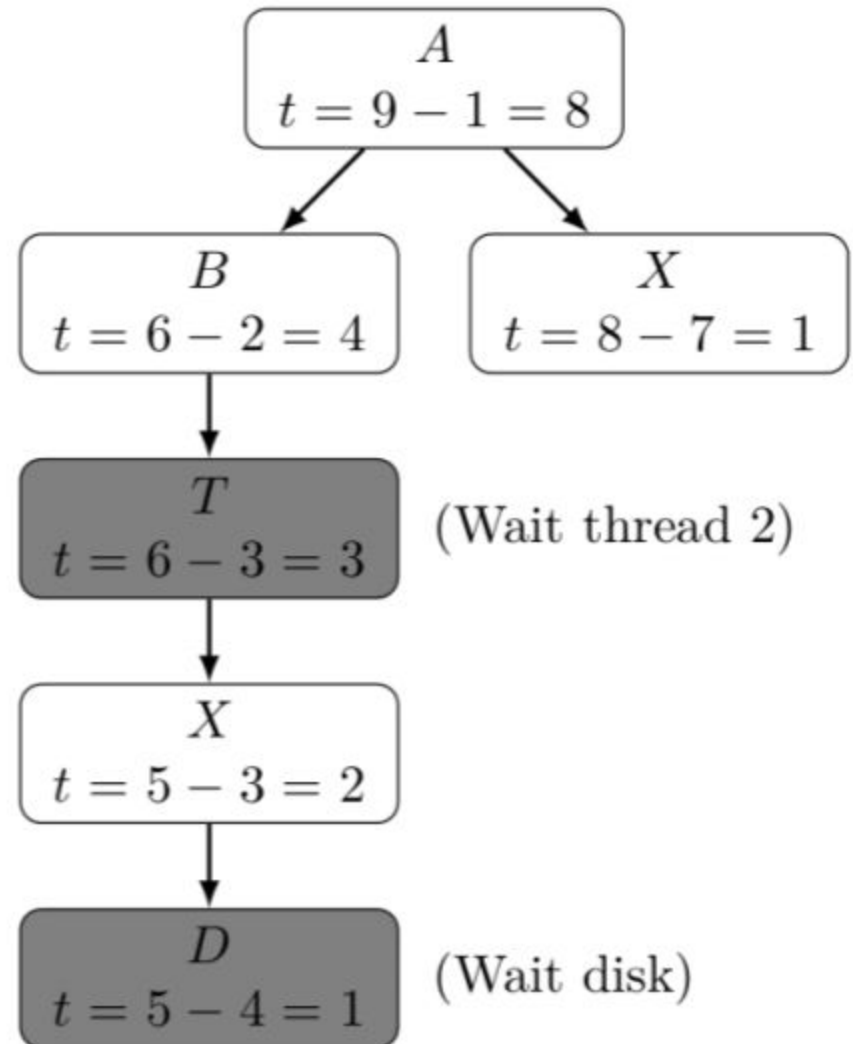
- Extract CCT for each request
- Extract critical path for each request
 - Graph of dependencies between threads



Merging CCT of Related Threads

28

Time	Thread 1	Thread 2
1	Call A	
2	Call B	
3	Wait thread 2	Call X
4		Wait disk
5		Return X
6	Return B	
7	Call X	
8	Return X	
9	Return A	



Comparing the Requests

29

- Comparing the requests ---> comparing the CCTs
- General solution:
 - Make critical CCT for each request
 - Cluster requests based on different features
 - Compare the groups to find the differences and reasons of slowness




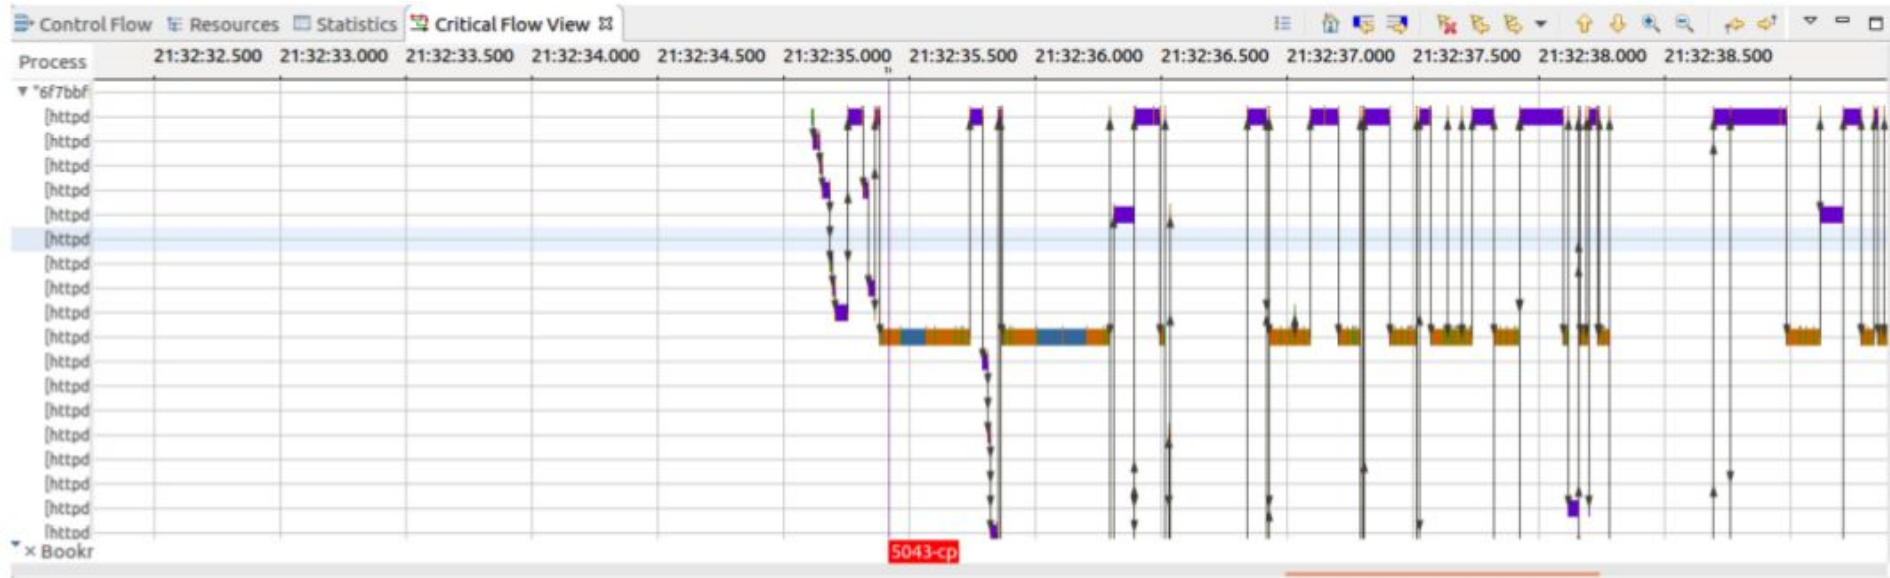
Use-cases

30

- Some period slow web requests
- One slow mysql request out of 10000 requests
- Real-time process was preempted with a low-priority thread doing some slow tasks
- Latencies related to:
 - CPU usage
 - Disk / network contention
 - Dependencies between threads
 - Futex/lock/etc.

Web Requests Periodic Slowdown

31



Performance

32

- How costly is this work?
 - Tracing
 - For every kernel event it adds 100 ns overhead
 - For the web application test it has 8.6 % overhead
 - Analysis
 - For 5 minutes of tracing (130 MB), it took 30 seconds to analyse and save the results and the size was 12.8 MB (10% of the original trace size)

Conclusion

33

- Automatic root cause identification of latency problems between multiple similar requests
 - Dynamic analysis through execution trace
 - Multilevel trace collection
 - Automatic grouping of requests based on their runtime behavior
 - Comparing different groups
- Support more latency problems
 - VMs
 - Containers
 - Distributed requests

References

- [1] Housseem Daoud, Naser Ezzati-Jivan, and Michel R. Dagenais. Dynamic trace-based sampling algorithm for memory usage tracking of enterprise applications. In 2017 IEEE High Performance Extreme Computing Conference (HPEC 2017), pages 1–7, Sept 2017. URL: <https://doi.org/10.1109/HPEC.2017.8091061>, doi:10.1109/HPEC.2017.8091061.
- [2] Simon Delise, Naser Ezzati-Jivan, and Michel R. Dagenais. Integrated modeling tool for indexing and analyzing state machine trace. submitted to Elsevier Journal of Computer and System Sciences, 2018.
- [5] Jean-Christian KOUAME, Naser Ezzati-Jivan, and Michel R. Dagenais. A data driven approach for pattern matching, filtering and analysis of execution traces. submitted to Springer Empirical Software Engineering, 2018.
- [3] K. Kouame, Naser Ezzati-Jivan, and Michel R. Dagenais. A flexible data-driven approach for execution trace filtering. In 2015 IEEE International Congress on Big Data, pages 698–703, June 2015. URL: <https://doi.org/10.1109/BigDataCongress.2015.112>, doi:10.1109/BigDataCongress.2015.112.
- [4] Alexander Montplaisir, Naser Ezzati-Jivan, Florian Wininger, and Michel R. Dagenais. State history tree: An incremental disk-based data structure for very large interval data. In 2013 International Conference on Social Computing, pages 716–724, Sept 2013. URL: <https://doi.org/10.1109/SocialCom.2013.107>, doi:10.1109/SocialCom.2013.107.
- [5] Loic Prieur-Drevon, Raphael Beamonte, Naser Ezzati-Jivan, and Michel R. Dagenais. Enhanced state history tree (esht): A stateful data structure for analysis of highly parallel system traces. In 2016 IEEE International Congress on Big Data (BigData Congress 2016), pages 83–90, June 2016. URL: <https://doi.org/10.1109/BigDataCongress.2016.19>, doi:10.1109/BigDataCongress.2016.19.
- [6] Fabien Reumont-Locke, Naser Ezzati-Jivan, and Michel R. Dagenais. Efficient methods for trace analysis parallelization. accepted in Springer International Journal of Parallel Programming, 2019.
- [7] Naser Ezzati-Jivan. Multi-Level Trace Abstraction, Linking and Display. PhD thesis, École Polytechnique de Montréal, 2014. URL: <http://publications.polymtl.ca/1402/>.
- [8] Doray, F. & Dagenais, M. R. (2016). Diagnosing performance variations by comparing multi-level execution traces. IEEE Transactions on Parallel and Distributed Systems, 28(2), p. 462-474. doi:10.1109/tpds.2016.2567390
- [9] Naser Ezzati-Jivan and Michel R. Dagenais. A stateful approach to generate synthetic events from kernel traces. Adv. Soft. Eng., 2012:6:6–6:6, January 2012. URL: <http://dx.doi.org/10.1155/2012/140368>, doi:10.1155/2012/140368.
- [10] Naser Ezzati-Jivan, Genevieve Bastien, and Michel R. Dagenais. High latency cause detection using multilevel dynamic analysis. In 2018 Annual IEEE International Systems Conference (SysCon 2018), Jan 2018.
- [11] Naser Ezzati-Jivan and Michel R. Dagenais. An efficient analysis approach for multi-core system tracing data. In Proceedings of the 16th International Conference on Software Engineering and Applications (SEA 2012), 2012. URL: <http://dx.doi.org/10.2316/P.2012.790-053>, doi:10.2316/P.2012.790-053.

Questions?

35

LTTng and Trace Compass

<http://lttng.org>

<http://tracecompass.org>

Sources:

<http://github.com/naser>

Contact:

ezzati@gmail.com



Silhouette Score

36

- Considers both cohesion and separation and is computed for a given sample RQ_i
 - $S_i = (\text{out } i - \text{In } i) / \max(\text{In } i, \text{out } i)$
 - In i : average distance of RQ_i to all other samples in its cluster
 - Out i : minimum average distance between RQ_i and out-cluster samples
 - The value of silhouette coefficient lies between -1 and 1, where a positive large value is more desirable.
 - We would like to have $\text{In } i$ as close as possible to zero for a very cohesive cluster, and to have $\text{out } i$ as large as possible representing very large separation from the closest neighbor cluster.

Silhouette Score

37

- while the quality of a clustering algorithm plays an important role in finding good clusters, having cohesive and well separated clusters also largely depends on the nature of the data and the distance metric.
- The silhouette coefficient can be averaged over a cluster or the entire samples to yield a silhouette score per cluster or for the entire clustering, respectively.
- We found that applying Kmeans to find three coarse clusters would yield the best total silhouette score of 0.42 and per cluster silhouette of $C0 = 0.38$, $C1 = 0.35$, $C2 = 0.48$.